


LoRA Fine-Tuning ELECTRA


comparing smaller vs. larger learning rate



Christopher Thompson

Last updated: May 20, 2026

✓ Section 1







☐ diff only**higher_LR****lower_LR**

> META (12 collapsed)

> CONFIG (201 collapsed)

✓ SUMMARY

> _wandb (1 collapsed)

✓ eval

accuracy	0.89905	0.86407
f1_weighted	0.89896	0.86335
loss	0.38796	0.51681
precision_weighted	0.90253	0.86972
recall_weighted	0.89905	0.86407
runtime	7.2839	7.6255
samples_per_second	274.715	262.408
steps_per_second	17.298	16.523
✓ train		
epoch	10	10
global_step	5010	5010
grad_norm	8.95733	19.47037
learning_rate	4.6719e-7	2.3360e-7
loss	0.24722	0.5569
_runtime	730	759
_step		
_timestamp	1709302566.80494	1709248415.93716
accuracy_gain	0.89155	0.85657
baseline_accuracy	0.0074963	0.0074963
baseline_f1_weighted	0.00085808	0.00085808

baseline_f1_weighted	0.896608	0.896608
test_accuracy	0.89643	0.85812
test_f1_weighted	0.89608	0.85706
total_flos	1662493908742080	1662493908742080
train_loss	0.88543	1.37232
train_runtime	704.4581	731.1902

Created with ❤️ on Weights & Biases.

<https://wandb.ai/chris-r-thompson1212-university-of-denver/llm-hw3/reports/LoRA-Fine-Tuning-ELECTRA---VmldzoxNjk1NTM4Ng>

